

SUR 2025/2026: Multimodal target-speaker detector

Dominik Huml (xhum1d00)

May 3, 2026

1. Task

The goal is a binary detector for the target speaker `m431`. Each evaluation sample is a face image and a paired voice recording sharing the same filename stem, so multimodal fusion at the sample level is straightforward. Three systems are required (image-only, audio-only, and a multimodal fusion of the two), each producing a real-valued score and a hard decision at prior 0.5. No pretrained models or outside data are allowed. Everything was trained from scratch on the provided 222-sample corpus with augmentation only.

2. Validation strategy

A plain K-Fold split leaks both session and speaker identity, since the target appears in only three sessions and most non-target speakers in one. I use a 3-fold leave-one-out scheme that groups target rows by session (01, 02, 03) and non-target rows by speaker, both fed to a single `GroupKFold` (`src/data/splits.py`). All three systems share the fold assignment, which lets the fusion stack on OOF scores. I pool the `_train` and `_dev` partitions into one 222-sample manifest before folding so that every target session takes a turn as held-out, rather than relying on the single train/dev split.

I report EER and min-DCF at prior 0.5, mean $\pm$ std across folds. For fusion I pool all 222 samples into a single OOF ranking and report the overall EER. Augmentation is train-only, val is always raw, and every fitted statistic (scaler, PCA basis, UBM, MAP target) is refit per fold. As a sanity check I ran a permutation test (E029). Shuffling train labels brought val EER back to 49 % (image) and 55 % (audio), confirming the models learn from labels and not auxiliary leakage.

3. Audio system

3.1 Front-end

I ablated MFCC, FBank, SDC, PLP and LPCC under identical UBM-MAP conditions. LPCC 13 + Δ + $\Delta\Delta$ at LPC order 12 won on both EER and min-DCF (4.21 % MFCC vs 3.33 % LPCC). Higher-dimensional alternatives over-parameterised GMM-32 on the available data, and order 12 had a 2.5 pp moat over its neighbours in the sweep.

Cepstral mean normalisation (CMN) is applied per utterance to normalise the channel response. CMVN regressed by 2 pp (E012) and was not adopted.

3.2 Classifier: UBM-MAP with tied covariance

The backbone is the standard UBM-MAP setup. A 32-component GMM is trained on all non-target training frames to get the UBM, then the target model is produced by MAP-adapting only the UBM means with relevance factor $\mathbf{r} = 16$. The score for an utterance is the per-frame log-likelihood ratio between the target and UBM models, averaged over all frames. I confirmed $\mathbf{r} \in \{4, 8, 16\}$ is a flat plateau on both diagonal and tied covariance (E013, E044) and rejected UBM-64 (E010, over-fit). The biggest single jump in the audio track came from changing the GMM covariance type (E037).

Diagonal underfits the off-diagonal correlations among LPCC coefficients. Full overfits at 48,672 parameters. Tied is the sweet spot. Numbers in Table 1.

3.3 Augmentation and TTA

Train-time augmentation applies two transforms to the LPCC pipeline. Pitch shift ± 1 –2 semitones (E025) was the key augmentation for LPCC. The same shift hurt MFCC in E014, so it was kept LPCC-specific. Codec simulation (E052) downsamples each utterance to 8 kHz and back, destroying frequencies above 4 kHz. The unaugmented system reached 0.46 % clean but collapsed to 13.33 % under bandwidth-limited val. Codec aug closes that gap at zero clean cost (Figure 1).

Covariance	EER (%)	min-DCF	Parameters
spherical	3.89 ± 3.75	0.0778	32
diagonal	4.35 ± 4.40	0.0870	1 248
tied	0.69 ± 0.98	0.0139	1 521
full	1.48 ± 0.92	0.0296	48 672

Table 1: GMM covariance ablation under fixed UBM-32, MAP $r = 16$, LPCC features. Tied is the empirical sweet spot.




Figure 1: Audio stress test (E051 + E052 paired). Codec aug closes the bandwidth failure mode ($13.33 \rightarrow 3.33$ %) and incidentally improves moderate-noise robustness with no clean-side regression.

At inference I apply speed-only test-time augmentation (E031). The utterance is scored at $1.0\times$, $0.9\times$ and $1.1\times$, and the three LLRs are averaged. Speed retiming preserves most of the spectral envelope, which is what LPCC encodes. Pitch shift at inference, on the other hand, collapses fold 0 to 9.86 %. Pitch is the right train-time augmentation but the wrong test-time one.

The final audio system is LPCC + tied-covariance UBM + MAP $r = 16$ + pitch and codec augmentation + speed TTA, reaching 0.46 ± 0.65 % EER at min-DCF 0.0093.

4. Image system

4.1 Features and classifier

The image classifier is 80×80 grayscale crops, per-pixel standardisation, 50-D PCA, logistic regression with $C = 1$. LBP (E005), Fisherfaces (E006), histogram equalisation and CLAHE (E041), and pyramid multi-scale PCA (E036) were tested and rejected. Sweeps confirmed $n_pca = 50$ (E011) and $C = 1$ (E040).

4.2 Augmentation: two-pass adversarial training (E033)

Augmentation is what made the image system competitive, and it runs in two passes.

Pass 1. Each training image is added in four versions: original, horizontal flip, brightness-scaled by $U[0.7, 1.3]$, and Gaussian pixel noise at $\sigma = 15$. A fresh PCA + LogReg is fit on the combined set. Brightness jitter is the key contributor in the ablation: session 03 has different lighting from sessions 01–02, and brightness scaling closes that gap.

Pass 2. For each training image I query the Pass-1 model at five angles in $[-10^\circ, +10^\circ]$ and pick the rotation that lands closest to the decision boundary. A rotated copy at that angle is added and PCA + LogReg is refit. Random rotation (E015) was less effective than adversarial selection in the ablation. Robustness numbers are summarised in Figure 2.

At inference I average the original and its horizontal flip. Rotation TTA was rejected (E030, eigenface corruption). E043 (flip + small-rotation TTA) was rejected after E049 failed to replicate the $+0.23$ pp gain.

The final image system reaches 0.51 ± 0.36 % EER at min-DCF 0.0102.




Figure 2: Image stress test (E051). AdvRot is robust to the photometric stresses tested (JPEG, blur, downsample stay at clean 0.51 %) and halves the geometric weaknesses. JPEG, blur, contrast, HE/CLAHE and Cutout 20×20 (E052) all regressed when added on top.

5. Fusion

Each modality produces a raw LLR (audio) or logit (image) on its own scale. I fit a Platt calibration per modality on its OOF scores (one-feature logreg, $C = 1e6$, balanced weights), then form a weighted sum on the 2-simplex ($w_{mfcc} + w_{lpcc} + w_{image} = 1$, non-negative) with weights chosen on a 51×51 grid that directly minimises OOF EER. Logistic-regression fusion was tested but underperformed grid search on this metric.

Fusion configuration		OOF EER (%)	min-DCF	Notes
MFCC + image	E009	3.75	0.0750	bimodal baseline
LPCC + image	E023	0.52	0.0104	LPCC replaces MFCC
MFCC + LPCC + image	E026/27	0.26	0.0052	trimodal, MFCC as tiebreaker
E052 + E033 + MFCC	E039	0.26	0.0052	new backbones, 0 / 222

Table 2: Score-level fusion progression. EER computed on the pooled OOF ranking of all 222 samples.

DET curves for the three streams plus the fused decision are plotted in Figure 3, and the disjoint-failure structure that makes the fusion possible is shown in Figure 4.

The grid converges to $w_{image} = 0.66$, $w_{lpcc} = 0.34$, $w_{mfcc} \approx 0$ (Table 2). The argmin is not unique on these scores: the simplex EER surface is flat at zero, and a re-run after recalibration produced $w_{image} = 0.54$, $w_{lpcc} = 0.46$ with the same 0/222 errors. Image carries the larger weight at the EER-optimal point on the simplex. LPCC adds complementary signal: 0 of 222 samples are misranked by both audio and image at their thresholds, so even though each modality misses a few targets the misses are disjoint. MFCC’s weight collapses to zero because MFCC and LPCC OOF scores are correlated at $r = 0.80$. I keep MFCC in the grid so that a future LPCC weakness can re-recruit it automatically. Quality-aware gating (E032), product-rule fusion (E046, E048) and score-level ensembles (E045) all regressed against simple weighted averaging.

6. Generalization and overfitting defences

I organise the safeguards into five risks, each paired with its defence.

Risk 1: capacity exceeds data size. Model sizes were chosen to match the data: UBM-32 (UBM-64 regressed, E010), PCA-50 (selected by sweep, E011), tied covariance with 1,521 parameters vs full at 48,672 (E037).

Risk 2: validation leaks session or speaker identity. Defence: Section 2 LOSO. Per-fold scaler, PCA, UBM and MAP target are refit, augmentation is train-only, and val is always raw.

Risk 3: hidden label leakage via auxiliary channels. Defence: permutation test (E029). Shuffling train labels brought val EER to 49 % and 55 %, both at chance. The 48 to 55 pp gap to the unshuffled flagships supports the conclusion that the models are not exploiting hidden label leakage.

Risk 4: post-hoc rationalisation across multi-axis changes. Defence: one knob per experiment, hypothesis




Figure 3: DET curves for the three modalities and the trimodal fusion. The fusion dot sits at the lower-left corner of the visible region. On this CV split it makes 0 errors out of 222 samples.

written before the run, failed runs kept in the log. Negative results (LBP, FBank, CMVN, HE/CLAHE, product-rule fusion) form part of the rationale.

Risk 5: distribution shift at eval time. Defence: stress testing (E028, E051, E052). The image flagship is robust to the photometric stresses tested: JPEG $q = 15$, blur $\sigma = 3$, and downsample 40 $\rightarrow$ 80 all stay at the clean 0.51 %. Remaining weaknesses are geometric (rot $> 15^\circ$: 7.6 %) and occlusion (11 %). Audio absorbs speed via TTA. Codec stress (13.33 % $\rightarrow$ 3.33 %) is what motivated E052. Additive noise raises audio EER to 4.35 % at 20 dB SNR and 6.85 % at 10 dB.

7. Results

The arc behind Table 3: MFCC + GMM baseline $\rightarrow$ UBM/MAP $\rightarrow$ LPCC $\rightarrow$ tied covariance $\rightarrow$ codec aug $\rightarrow$ adversarial rotation on the image side $\rightarrow$ trimodal score-level fusion. Figure 5 plots the same trajectory across all 52 experiments.

Reproduction

All systems run from a clean clone with

```
uv sync
uv run predict_audio.py --eval-dir <dir> --output results/audio_lpcc_tied_codecaug.txt
uv run predict_image.py --eval-dir <dir> --output results/image_pca_adv_rot.txt
uv run predict_fusion.py --eval-dir <dir> --output results/fusion_trimodal.txt
```

Each output file contains one line per evaluation sample with three whitespace-separated fields: stem, score (higher = more confident target), and a hard decision (1 = target, 0 = non-target) thresholded at the Bayes optimum for prior 0.5, calibrated on OOF min-DCF.

I submit three result files, one per row of the flagship block in Table 3: `audio_lpcc_tied_codecaug.txt` (E052), `image_pca_adv_rot.txt` (E033), and `fusion_trimodal.txt` (E039). The full ablation arc is documented in this report rather than reproduced as separate result files.




Figure 4: Audio $\times$ image complementarity at per-stream EER thresholds. 15 samples are rescued by audio alone, 5 by image alone, none are missed by both. This disjoint failure pattern is what makes weighted-sum fusion reach 0 errors.

System		EER (%)	min-DCF
<i>Audio</i>			
MFCC + GMM	E001	17.92 ± 7.81	0.2250
MFCC + UBM/MAP + aug	E008	4.21 ± 3.11	0.0509
LPCC + tied covariance	E037	0.69 ± 0.98	0.0139
LPCC + tied + pitch & codec + speed TTA	E052	0.46 ± 0.65	0.0093
<i>Image</i>			
PCA + LogReg	E004	4.49 ± 4.26	0.0565
PCA + LogReg + aug	E007	0.97 ± 0.86	0.0194
PCA + LogReg + aug + AdvRot	E033	0.51 ± 0.36	0.0102
<i>Fusion</i> [†]			
MFCC + image (bimodal baseline)	E009	3.75	0.0750
LPCC + image	E023	0.52	0.0104
MFCC + LPCC + image (E008 + E007 backbones)	E027	0.26	0.0052
Trimodal (E052 + E033 + MFCC, new backbones)	E039	0.26	0.0052

Table 3: Final cross-validation results across the three modalities and the trimodal fusion. [†] Fusion EER is OOF on the full 222-sample pool, where the flagship E039 makes 0 / 222 errors. Audio and image rows are mean $\pm$ std across the 3 LOSO folds.



Figure 5: Project arc across 52 experiments. Y-axis is EER on a log scale. Each modality flagship reduced EER substantially: audio 97 %, image 89 %, fusion 93 % relative EER reduction.